

Introduction

Every automation platform eventually forces you to answer the same question: what happens when an item can go more than two places? In n8n, the naive answer is to nest IF nodes — check one condition, and inside the "false" branch, check another, and inside that one, another. It works for a while. Then a new business rule shows up, and the workflow that used to fit on one screen now needs a scrollbar just to follow the logic.

This isn't a n8n-specific problem. It's the same shape as an `if / elif / elif / else` chain in Python, or a stack of nested ternaries in JavaScript — anywhere a developer has to route one input to one of several outcomes. If you've written code, you've already solved this problem once with a `switch` statement or a dictionary dispatch. n8n's Switch node is that same idea, expressed as a node instead of a keyword.

This lesson matters for anyone building real automations, not just toy examples. Support ticket routing, order processing, lead qualification, alert triage — almost every practical workflow has more than two outcomes. Knowing when to reach for Switch instead of nested IFs is the difference between a workflow you can debug in thirty seconds and one you're afraid to touch.

Problem Overview

Picture a support ticket workflow. Tickets need to go to one of four places depending on their type: Billing, Technical, Refund, or Complaint. Using only IF nodes, you'd write it like this:

- IF ticket type is Billing → route to Billing
- Else, IF ticket type is Technical → route to Technical
- Else, IF ticket type is Refund → route to Refund
 - Else, route to Complaint

Each new category means one more IF node nested inside the previous one's "false" branch. The workflow grows sideways and downward at the same time. It's not just visually messy — it's cognitively expensive. To understand what happens to a "Refund" ticket, you have to trace through every IF node that comes before it, confirming each one evaluates to false.

The root cause is structural: an IF node is a binary decision. It was designed to answer exactly one yes/no question and hand the item to one of two outputs. The moment your real-world logic has three or more distinct outcomes, you're bending a binary tool into a multi-way shape it was never built for.

Example

Take four incoming tickets:

Ticket Text	Ticket Type	Field Expected Output
"I was charged twice this month"	billing	Billing
"The app crashes on login"	technical	Technical
"I want my money back"	refund	Refund
"Your support team ignored me"	complaint	Complaint

With nested IFs, ticket 4 (Complaint) has to fail three separate checks before it lands anywhere — even though "complaint" is a completely known, expected category. It's being treated as leftover logic rather than a first-class branch, purely because of how the tool forces you to express the logic.

With a Switch node, each of these four types gets its own named output, evaluated independently against the same field. No ticket is a "leftover" — every branch is equally visible.

Intuition

Here's how an experienced automation engineer actually notices this problem: they stop looking at the individual IF nodes and start looking at the shape of the whole workflow. If a workflow's nodes cascade diagonally down the canvas, each one nested inside the last, that's the visual signature of a decision that has too many outcomes for the tool being used.

The next question is: how many *distinct*, mutually exclusive outcomes are actually being tested? If a developer were writing this in code, they'd immediately reach for one of two shapes:

```
if ticket_type == "billing":
    ...
elif ticket_type == "technical":
    ...
elif ticket_type == "refund":
    ...
else:
    ...
```

That's a linear `elif` chain — semantically identical to nested IFs, just written flatter. And it's exactly the shape a `switch` statement (or n8n's Switch node) exists to replace. The insight isn't "IF is bad" — IF is perfect for a true two-way fork. The insight is that **the number of real branches should determine the node**, not habit. Two outcomes: reach for IF, it's simpler. Three or more real, mutually exclusive outcomes: reach for Switch, because it turns a nested tree into a flat list of named cases.

Brute Force Solution

The "brute force" approach here is the nested-IF pyramid described above.

Idea: Chain IF nodes, testing one condition per node, pushing untested items down to the next check.


Algorithm: 1. IF node 1 tests condition A. True → output A. False → continue. 2. IF node 2 (nested in node 1's false branch) tests condition B. True → output B. False → continue. 3. Repeat for each additional condition. 4. The final "false" branch becomes the catch-all.

Advantages: - Works with a node every n8n user already knows. - Fine, even preferable, when there are genuinely only two outcomes.

Disadvantages: - Grows linearly in nesting depth as branches are added — a 5-branch decision needs 4 nested IF nodes. - Readability collapses fast; you have to mentally track "false, false, false" to know what a deep branch means. - Reordering or inserting a new condition means restructuring the nesting, not just adding a node. - Time and space complexity in execution are trivial ($O(n)$ condition checks for n branches, negligible memory) — the cost here isn't runtime, it's human comprehension and maintenance.

Optimal Solution

The Switch node flattens the tree into a list. Instead of nesting, you attach one rule per output, and all outputs sit side by side on one node.



Every rule in a Switch node is built from three pieces:

Piece	Role
Field	The property being read off the incoming item (e.g., <code>ticket_type</code>)
Condition	What that field must satisfy (e.g., equals <code>"refund"</code>)
Output	The named path the item takes when the condition matches

Two behaviors matter more than anything else here:

1. **Rules evaluate top to bottom, and it's first-match-wins.** If a ticket matches both a "refund" rule and a "billing" rule, and refund is listed first, it exits through Refund and never reaches the billing check — even though it would also satisfy it. This is the same semantics as a `switch` statement with implicit breaks, not a set of independent, always-fires-when-true conditions. Reordering rules changes routing outcomes, which makes rule order a design decision, not an implementation detail.
2. **There's a built-in fallback output.** Any item that matches none of the defined rules still needs somewhere to go — the fallback output exists specifically so nothing silently disappears from the workflow.

Python Code

While Switch is an n8n node, the underlying logic is worth writing in code, because it's exactly what you'd reach for in any backend service doing the same routing:

```
from typing import Callable

def route_ticket(ticket_type: str, rules: list[tuple[str, str]], fallback: str = "unrouted") -> str:
    """
    First-match-wins router, mirroring n8n's Switch node semantics.

    rules: ordered list of (expected_value, output_label) pairs.
    Returns the output label for the first matching rule, or `fallback`.
    """
    for expected_value, output_label in rules:
        if ticket_type == expected_value:
            return output_label
    return fallback

# Rule order matters: refund is checked before billing
ticket_rules = [
    ("refund", "Refund"),
    ("billing", "Billing"),
    ("technical", "Technical"),
    ("complaint", "Complaint"),
]

print(route_ticket("refund", ticket_rules))    # Refund
print(route_ticket("unknown", ticket_rules))  # unrouted
```

Code Walkthrough

- `rules` is a list of tuples, not a dict, because **order matters** — the exact behavior we need to mirror from Switch. A dict would lose the ordering guarantee (or at least obscure it).
- The function iterates rules in order and returns as soon as one matches — this is the "first-match-wins" behavior, directly equivalent to Switch evaluating rules top to bottom.
- If the loop completes without a match, `fallback` is returned — this is the Switch node's fallback output, made explicit instead of letting the item vanish.
- Each tuple unpacks into `(expected_value, output_label)` — this maps directly onto a Switch rule's "condition" and "output" pieces; the "field" is fixed here as `ticket_type` but could be parameterized with a `Callable` extractor for a more general router.

Dry Run

Input: a ticket with `ticket_type = "refund"`, using `ticket_rules` above.

1. Loop starts. First rule: `("refund", "Refund")`.
2. Check: does `"refund" == "refund"` ? Yes.
3. Return `"Refund"` immediately — the loop never reaches the `"billing"` rule, even though nothing here says the ticket *couldn't* also match billing logic elsewhere.

Now input: `ticket_type = "spam"`.

1. Check against `"refund"` — no match.
2. Check against `"billing"` — no match.
3. Check against `"technical"` — no match.
4. Check against `"complaint"` — no match.
5. Loop exhausted. Return `fallback`, `"unrouted"`.

This is exactly how a ticket with no matching rule in n8n's Switch node lands in the fallback output instead of disappearing from the workflow.

Complexity Analysis

- **Time:** $O(n)$ where n is the number of rules, since in the worst case (no match, or the last rule matches) every rule is checked once. This is correct because there's no way to know which rule will match without checking it — there's no shortcut without additional structure (like a lookup table if all rules are pure equality checks, discussed below).
- **Space:** $O(1)$ beyond the input, since routing doesn't build any auxiliary structure — it's a simple linear scan.

- This mirrors the Switch node's actual execution cost: negligible per item, since rule counts in real workflows are small (single digits). The complexity concern was never about runtime — it's about human readability, and Switch fixes that without changing the algorithmic shape.

Alternative Solutions

If every rule is a simple equality check on the same field, a dictionary lookup is a reasonable alternative:

```
def route_ticket_dict(ticket_type: str, mapping: dict[str, str], fallback: str = "unrouted") -> str:
    return mapping.get(ticket_type, fallback)
```

This is $O(1)$ average-case instead of $O(n)$, and it's tempting to call it "more optimal." But it only works when rules are exact-match and mutually exclusive by construction — it silently can't express "refund OR billing" or "amount > 500," and it loses explicit ordering semantics (Python dicts preserve insertion order, but nothing about the lookup *depends* on it, so a future maintainer could reorder keys without realizing it changes behavior). Switch's rule-based approach trades a small performance cost for expressiveness and explicit ordering — the right trade for automation workflows where clarity beats microseconds.

Edge Cases

- **No ticket type field present at all:** the field read comes back empty/null; make sure every rule's condition handles a missing value explicitly rather than assuming it will simply fail every check silently.
- **Duplicate-matching item** (matches two rules): resolved entirely by rule order — the first listed rule wins, so intentional ordering is required whenever categories can overlap.
- **Empty rule set:** every item goes straight to fallback; useful to confirm fallback is actually wired to something, not left dangling.
- **Only one real rule defined:** functionally identical to an IF node — a sign Switch may be overkill for this specific decision.
- **Case-sensitive field values** ("Refund" vs "refund"): a mismatch here silently sends items to fallback instead of raising an error, so normalizing case before the Switch node avoids "phantom" fallback items.

Common Mistakes

1. **Assuming multiple rules can all fire for one item.** Switch is first-match-wins; only one output receives each item, unlike systems where every matching condition triggers independently.
2. **Ignoring rule order.** Rules are evaluated top to bottom; two correct-looking rules can produce wrong routing purely because of their position.
3. **Forgetting to wire the fallback output.** Items that match nothing don't error — they just disappear if fallback isn't connected to anything.
4. **Reaching for Switch when there are only two real outcomes.** It adds setup for no benefit over a plain IF node.
5. **Testing the wrong field per rule.** Each rule reads a field independently — a typo or wrong field reference in one rule fails silently rather than throwing an obvious error.
6. **Not accounting for type mismatches** (comparing a string field against a number condition), which usually evaluates to false rather than raising a visible error.

Interview Questions

- How would you design a router that dispatches based on more than two conditions? (Expects: switch/dispatch-table reasoning, not nested conditionals.)
- What's the time complexity difference between a first-match-wins linear scan and a dictionary-based dispatch, and when does that difference actually matter?
- How do you handle an input that matches none of your defined cases?
- What's the tradeoff between an `elif` chain and a dictionary dispatch table in Python?
- How would you detect, in code review, that a decision structure has outgrown a binary IF/else pattern?

Similar Problems

- **LeetCode 1046 (Last Stone Weight)** — not routing-related directly, but a good exercise in choosing the right data structure (heap) once a naive approach (sorting repeatedly) becomes the bottleneck, the same "right tool for the shape of the problem" instinct.
- **Implementing a state machine** — any problem involving discrete states with more than two transitions (e.g., LeetCode-style traffic light or vending machine simulations) uses the identical first-match/dispatch-table logic.
- **String pattern dispatch problems** (e.g., command parsers) — routing a command string to one of several handlers is structurally the same problem as ticket routing.

Key Takeaways

- An IF node (or an `if/else` statement) is built for exactly two outcomes — forcing it to handle three or more just nests it repeatedly.
- A Switch node (or a `switch` /dispatch-table pattern in code) flattens multi-way decisions into a single, readable list of named branches.
- Switch is first-match-wins: rule order directly determines routing, and reordering rules changes behavior.
- Always wire a fallback path — items matching nothing should land somewhere explicit, never disappear.
- Reach for the simpler tool when the problem is actually simple: two outcomes still means IF, not Switch.

Watch the Video

For the full walkthrough with the ticket-routing workflow built live, watch the video here: {LINK}

About the Series

This lesson is part of **Fun with Learning Technology**, a series built for developers who want practical, no-fluff explanations of the tools shaping modern software — from core algorithms and interview prep to automation platforms like n8n. Each lesson takes one concept, shows why it matters, and builds the intuition behind it rather than just the steps.

Call To Action

If this helped clarify when to reach for Switch instead of nesting IF nodes, give the video a like on YouTube and subscribe for the rest of the n8n lessons series. For deeper write-ups like this one delivered straight to your inbox, subscribe to the Substack. Drop a comment with your own nested-IF horror story, and share this with a teammate who's still three IF nodes deep in a workflow they're afraid to touch.

Quick Review

n8n: when do you reach for Switch over IF?

When items route into 3+ mutually exclusive paths by rule, not just two — IF only has true/false.

Switch node: how many outputs can a single item hit?

Exactly one — first matching rule wins, top to bottom; item doesn't fan out to multiple branches.

Why avoid nesting IFs for multi-way routing?

Nested IFs form a pyramid of conditions that's hard to read/maintain; Switch flattens it into one node.

Trickiest Switch bug: items matching no rule?

They fall through silently unless you wire the fallback/default output — always catch it explicitly.

What does a Switch rule actually evaluate?

A condition (expression or value comparison) checked per item; first true rule sends that item to its output.

Interview follow-up: two Switch rules both match an item — which wins?

The first one in top-to-bottom order; later matching rules are never evaluated for that item.